

Rules for every challenge

- Do all setup (filling, seeding indices, printing) outside the timed block.
- Put only the target work between t_0 and t_1 .
- Repeat ≥ 5 trials and average. Try multiple n (e.g., $1e3$, $1e4$, $1e5$) where it makes sense.
- Hand in: your code, a small table of times, and a 2–4 sentence explanation per part.

Part A – Circular & Doubly Linked List mastery

A1 – CSLL: tail-to-head wrap vs manual reset

- Implement traversal of n nodes in two ways:
 - (i) CSLL with $\text{tail} \rightarrow \text{next} = \text{head}$ and a single loop of n steps;
 - (ii) Non-circular SLL that restarts at head whenever you hit nullptr.
- Predict which is faster and why (branching, cache/predictability).
- Measure and explain.

A2 – CSLL deletion with/without predecessor

- **Case 1:** Delete a given node when you also have its predecessor ($O(1)$).
- **Case 2:** Delete the same node when you only have the node pointer (must find predecessor).
- Predict cost difference; measure for random positions; explain the curve.

A3 – Rotate- k on CSLL vs SLL

- Implement “rotate right by k ”: in CSLL it’s pointer moves; in SLL it’s find-break-relink.
- Test for multiple k (small, $n/2$, $n-1$).
- Decide which wins and under what k ranges.

A4 – DLL vs SLL: erase-given-node

- Build DLL with prev. Given a pointer to any node, erase it.
- Compare to SLL erase with known predecessor and SLL erase without predecessor.
- Predict \rightarrow measure \rightarrow explain $O(1)$ vs $O(n)$ and constant-factor hits.

A5 – Push/pop ends: head-only vs head+tail

- Implement push_front/pop_front and push_back/pop_back on:
 - (i) SLL with head only, (ii) SLL with head+tail, (iii) DLL with head+tail.
- Benchmark each op for random mixes (e.g., 70% push_back, 30% pop_front).
- Explain why tail changes the story.

A6 – Memory overhead audit

- For the same logical dataset size n , allocate SLL, CSLL, and DLL nodes.
- Report bytes per node (pointer count), total bytes, and measured allocation time.
- Discuss the time–space trade-off you’d choose for frequent middle deletions.

Part B – Real-world use cases

B1 – Recent Items Tray (add/remove at the same end)

- Build a “recent items” tray where the most recently added item is always the next one removed.

- **Implement with:**
 - A) Singly linked nodes adding/removing at the front.
 - B) Doubly linked nodes adding/removing at the front.
- **Predict which is faster** (pointer count vs rewiring), measure adds/removes (mixed workload), explain whether the second pointer helps here.

B2 — Editor Undo History

- **Implement an Undo history** where each new action is “placed on top,” and Undo removes the last action added.
- **Version A: Singly linked front-add/front-remove.**
- **Version B: Dynamic array (grow by doubling).**
- **Workload: 80% add actions, 20% undo actions.**
- **Predict throughput and memory spikes (reallocs) vs constant-time links; measure and justify which design you’d ship.**